

```
In [3]: #import packages and clean data before running multiple regression analysis. Rename th
import numpy as np
import pandas as pd
from sklearn import linear_model
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
pd.set_option('display.max_columns', None)
import pylab
from pylab import rcParams
import statsmodels.api as sm
import statistics
from scipy import stats
import sklearn
from sklearn import preprocessing
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.metrics import classification_report
from scipy.stats import chi-square
from scipy.stats import chi2_contingency
df = pd.read_csv (r'C:\Users\fahim\Documents\0_WGUDocuments\d208\1medical_clean.csv')
df.rename(columns={'Item1':'Timely_admis','Item2':'Timely_treat',
'Item3':'Timely_visits','Item4':'Reliability',
'Item5':'Options','Item6':'Hrs_treat',
'Item7':'Courteous','Item8':'Active_listen'},inplace=True)
df.head()
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 50 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CaseOrder              10000 non-null  int64
1   Customer_id            10000 non-null  object
2   Interaction             10000 non-null  object
3   UID                    10000 non-null  object
4   City                   10000 non-null  object
5   State                  10000 non-null  object
6   County                 10000 non-null  object
7   Zip                    10000 non-null  int64
8   Lat                    10000 non-null  float64
9   Lng                    10000 non-null  float64
10  Population              10000 non-null  int64
11  Area                    10000 non-null  object
12  TimeZone                10000 non-null  object
13  Job                     10000 non-null  object
14  Children                10000 non-null  int64
15  Age                     10000 non-null  int64
16  Income                  10000 non-null  float64
17  Marital                 10000 non-null  object
18  Gender                  10000 non-null  object
19  ReAdmis                 10000 non-null  object
20  VitD_levels             10000 non-null  float64
21  Doc_visits              10000 non-null  int64
22  Full_meals_eaten        10000 non-null  int64
23  vitD_supp               10000 non-null  int64
24  Soft_drink              10000 non-null  object
25  Initial_admin            10000 non-null  object
26  HighBlood               10000 non-null  object
27  Stroke                  10000 non-null  object
28  Complication_risk        10000 non-null  object
29  Overweight              10000 non-null  object
30  Arthritis               10000 non-null  object
31  Diabetes                10000 non-null  object
32  Hyperlipidemia          10000 non-null  object
33  BackPain                10000 non-null  object
34  Anxiety                 10000 non-null  object
35  Allergic_rhinitis        10000 non-null  object
36  Reflux_esophagitis       10000 non-null  object
37  Asthma                  10000 non-null  object
38  Services                 10000 non-null  object
39  Initial_days             10000 non-null  float64
40  TotalCharge              10000 non-null  float64
41  Additional_charges       10000 non-null  float64
42  Timely_admis             10000 non-null  int64
43  Timely_treat             10000 non-null  int64
44  Timely_visits            10000 non-null  int64
45  Reliability              10000 non-null  int64
46  Options                  10000 non-null  int64
47  Hrs_treat                10000 non-null  int64
48  Courteous                10000 non-null  int64
49  Active_listen            10000 non-null  int64
dtypes: float64(7), int64(16), object(27)
memory usage: 3.8+ MB

```

```

In [4]: #check if there are any missing data entries - if there are none then the output should
df.isna().any()

```

```
Out[4]: CaseOrder      False
Customer_id    False
Interaction     False
UID            False
City           False
State          False
County         False
Zip            False
Lat            False
Lng            False
Population     False
Area           False
TimeZone       False
Job            False
Children       False
Age            False
Income         False
Marital        False
Gender         False
ReAdmis        False
VitD_levels    False
Doc_visits     False
Full_meals_eaten False
vitD_supp      False
Soft_drink     False
Initial_admin  False
HighBlood      False
Stroke         False
Complication_risk False
Overweight     False
Arthritis      False
Diabetes       False
Hyperlipidemia False
BackPain       False
Anxiety        False
Allergic_rhinitis False
Reflux_esophagitis False
Asthma         False
Services       False
Initial_days   False
TotalCharge    False
Additional_charges False
Timely_admis   False
Timely_treat   False
Timely_visits  False
Reliability    False
Options        False
Hrs_treat      False
Courteous      False
Active_listen  False
dtype: bool
```

```
In [5]: #check if there is any duplicate data entries present in columns
df[df.duplicated()]
```

```
Out[5]: CaseOrder  Customer_id  Interaction  UID  City  State  County  Zip  Lat  Lng  Population  Area
```

```
In [6]: # check if there are any duplicated columns in the data set - if there are none then t
df.columns.duplicated().any()
```

```
Out[6]: False
```

```
In [7]: # check if there are any duplicated rows in the data set - if there are none then the
df.duplicated().any()
```

```
Out[7]: False
```

```
In [10]: # remove demographic data from the data set since these entries won't be necessary for
df = df.drop(['CaseOrder', 'Customer_id', 'Interaction', 'UID', 'City', 'State', 'County', 'Z
```

```
In [11]: # check to make sure that the columns for demographic data were dropped before proceed
df.head()
```

```
Out[11]:
```


	Children	Age	Income	Marital	Gender	ReAdmis	VitD_levels	Doc_visits	Full_meals_eaten	vi
0	1	53	86575.93	Divorced	Male	No	19.141466	6	0	
1	3	51	46805.99	Married	Female	No	18.940352	4	2	
2	3	53	14370.14	Widowed	Female	No	18.057507	4	1	
3	0	78	39741.49	Married	Male	No	16.576858	4	1	
4	1	22	1209.56	Widowed	Female	No	17.439069	5	0	

```
In [12]: # convert categorical yes/no values to numeric 1/0 values
df = df.replace(to_replace = ['Yes', 'No'], value = [1,0])
df
```

Out[12]:

	Children	Age	Income	Marital	Gender	ReAdmis	VitD_levels	Doc_visits	Full_meals_eaten
0	1	53	86575.93	Divorced	Male	0	19.141466	6	0
1	3	51	46805.99	Married	Female	0	18.940352	4	2
2	3	53	14370.14	Widowed	Female	0	18.057507	4	1
3	0	78	39741.49	Married	Male	0	16.576858	4	1
4	1	22	1209.56	Widowed	Female	0	17.439069	5	0
...	...	...	...	...	...	...	...	...	...
9995	2	25	45967.61	Widowed	Male	0	16.980860	4	2
9996	4	87	14983.02	Widowed	Male	1	18.177020	5	0
9997	3	45	65917.81	Separated	Female	1	17.129070	4	2
9998	3	43	29702.32	Divorced	Male	1	19.910430	5	2
9999	8	70	62682.63	Separated	Female	1	18.388620	5	0

10000 rows × 36 columns



In [14]:

```
# convert the categorical variable of genders to a numeric variable
df['Gender'] = df['Gender'].replace(['Male', 'Female', 'Nonbinary'],[1,2,3])
df
```

Out[14]:

	Children	Age	Income	Marital	Gender	ReAdmis	VitD_levels	Doc_visits	Full_meals_eaten
<b>0</b>	1	53	86575.93	Divorced	1	0	19.141466	6	0
<b>1</b>	3	51	46805.99	Married	2	0	18.940352	4	2
<b>2</b>	3	53	14370.14	Widowed	2	0	18.057507	4	1
<b>3</b>	0	78	39741.49	Married	1	0	16.576858	4	1
<b>4</b>	1	22	1209.56	Widowed	2	0	17.439069	5	0
...	...	...	...	...	...	...	...	...	...
<b>9995</b>	2	25	45967.61	Widowed	1	0	16.980860	4	2
<b>9996</b>	4	87	14983.02	Widowed	1	1	18.177020	5	0
<b>9997</b>	3	45	65917.81	Separated	2	1	17.129070	4	2
<b>9998</b>	3	43	29702.32	Divorced	1	1	19.910430	5	2
<b>9999</b>	8	70	62682.63	Separated	2	1	18.388620	5	0

10000 rows × 36 columns

```
In [15]: # convert the non-married Marital status values to "Married/Not Married", then convert
#this will make the Marital variable easier to work with during regression analysis
df['Marital'] = df['Marital'].replace(['Divorced','Widowed','Separated','Never Married'],[1,0])
df['Marital'] = df['Marital'].replace(['Married','Not Married'],[1,0])
```

```
In [32]: # convert the Initial_Admin, Complication_risk, and Services variables into integers to
df['Initial_admin'] = df['Initial_admin'].replace(['Elective Admission','Observation Admission'],[1,0])
df['Complication_risk'] = df['Complication_risk'].replace(['Low','Medium','High'],[1,2,3])
df['Services'] = df['Services'].replace(['Blood Work','CT Scan','Intravenous','MRI'],[1,2,3,4])
df.info()
df.describe()
my_list = df.columns.values.tolist()
print(my_list)
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 10000 entries, 0 to 9999
```

```
Data columns (total 37 columns):
```

#	Column	Non-Null Count	Dtype
0	Initial_days	10000 non-null	float64
1	Children	10000 non-null	int64
2	Age	10000 non-null	int64
3	Income	10000 non-null	float64
4	Marital	10000 non-null	int64
5	Gender	10000 non-null	int64
6	ReAdmis	10000 non-null	int64
7	VitD_levels	10000 non-null	float64
8	Doc_visits	10000 non-null	int64
9	Full_meals_eaten	10000 non-null	int64
10	vitD_supp	10000 non-null	int64
11	Soft_drink	10000 non-null	int64
12	Initial_admin	10000 non-null	int64
13	HighBlood	10000 non-null	int64
14	Stroke	10000 non-null	int64
15	Complication_risk	10000 non-null	int64
16	Overweight	10000 non-null	int64
17	Arthritis	10000 non-null	int64
18	Diabetes	10000 non-null	int64
19	Hyperlipidemia	10000 non-null	int64
20	BackPain	10000 non-null	int64
21	Anxiety	10000 non-null	int64
22	Allergic_rhinitis	10000 non-null	int64
23	Reflux_esophagitis	10000 non-null	int64
24	Asthma	10000 non-null	int64
25	Services	10000 non-null	int64
26	TotalCharge	10000 non-null	float64
27	Additional_charges	10000 non-null	float64
28	Timely_admis	10000 non-null	int64
29	Timely_treat	10000 non-null	int64
30	Timely_visits	10000 non-null	int64
31	Reliability	10000 non-null	int64
32	Options	10000 non-null	int64
33	Hrs_treat	10000 non-null	int64
34	Courteous	10000 non-null	int64
35	Active_listen	10000 non-null	int64
36	intercept	10000 non-null	int64

```
dtypes: float64(5), int64(32)
```

```
memory usage: 2.8 MB
```

```
['Initial_days', 'Children', 'Age', 'Income', 'Marital', 'Gender', 'ReAdmis', 'VitD_levels', 'Doc_visits', 'Full_meals_eaten', 'vitD_supp', 'Soft_drink', 'Initial_admin', 'HighBlood', 'Stroke', 'Complication_risk', 'Overweight', 'Arthritis', 'Diabetes', 'Hyperlipidemia', 'BackPain', 'Anxiety', 'Allergic_rhinitis', 'Reflux_esophagitis', 'Asthma', 'Services', 'TotalCharge', 'Additional_charges', 'Timely_admis', 'Timely_treat', 'Timely_visits', 'Reliability', 'Options', 'Hrs_treat', 'Courteous', 'Active_listen', 'intercept']
```

```
C:\Users\fahim\AppData\Local\Temp\ipykernel_21320\1287141306.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df['Initial_admin'] = df['Initial_admin'].replace(['Elective Admission','Observation Admission','Emergency Admission'],[1,2,3])
C:\Users\fahim\AppData\Local\Temp\ipykernel_21320\1287141306.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df['Complication_risk'] = df['Complication_risk'].replace(['Low','Medium','High'],[1,2,3])
C:\Users\fahim\AppData\Local\Temp\ipykernel_21320\1287141306.py:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df['Services'] = df['Services'].replace(['Blood Work','CT Scan','Intravenous','MRI'],[1,2,3,4])
```

Out[32]:

	Initial_days	Children	Age	Income	Marital	Gender	ReAdmis	VitD_levels	Doc_visits	Full_n
0	10.585770	1	53	86575.93	0	1	0	19.141466	6	
1	15.129562	3	51	46805.99	1	2	0	18.940352	4	
2	4.772177	3	53	14370.14	0	2	0	18.057507	4	
3	1.714879	0	78	39741.49	1	1	0	16.576858	4	
4	1.254807	1	22	1209.56	0	2	0	17.439069	5	
...	...	...	...	...	...	...	...	...	...	...
9995	51.561220	2	25	45967.61	0	1	0	16.980860	4	
9996	68.668240	4	87	14983.02	0	1	1	18.177020	5	
9997	70.154180	3	45	65917.81	0	2	1	17.129070	4	
9998	63.356900	3	43	29702.32	0	1	1	19.910430	5	
9999	70.850590	8	70	62682.63	0	2	1	18.388620	5	

10000 rows × 37 columns

```
In [37]: # move the chosen target variable "Initial_days" to beginning of the columns
df=df[['Initial_days','Children', 'Age', 'Income', 'Marital', 'Gender', 'ReAdmis', 'Vi
# Confirm that the target variable was moved before exporting the prepared dataset
my_list = df.columns.values.tolist()
print(my_list)
# describe the dataframe to identify distribution of variables
df.describe()
```



```
['Initial_days', 'Children', 'Age', 'Income', 'Marital', 'Gender', 'ReAdmis', 'VitD_levels', 'Doc_visits', 'Full_meals_eaten', 'vitD_supp', 'Soft_drink', 'Initial_admin', 'HighBlood', 'Stroke', 'Complication_risk', 'Overweight', 'Arthritis', 'Diabetes', 'Hyperlipidemia', 'BackPain', 'Anxiety', 'Allergic_rhinitis', 'Reflux_esophagitis', 'Asthma', 'Services', 'TotalCharge', 'Additional_charges', 'Timely_admis', 'Timely_treat', 'Timely_visits', 'Reliability', 'Options', 'Hrs_treat', 'Courteous', 'Active_listen']
```

Out[37]:

	Initial_days	Children	Age	Income	Marital	Gender	ReA
<b>count</b>	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
<b>mean</b>	34.455299	2.097200	53.511700	40490.495160	0.202300	1.544600	0.360000
<b>std</b>	26.309341	2.163659	20.638538	28521.153293	0.401735	0.539296	0.480000
<b>min</b>	1.001981	0.000000	18.000000	154.080000	0.000000	1.000000	0.000000
<b>25%</b>	7.896215	0.000000	36.000000	19598.775000	0.000000	1.000000	0.000000
<b>50%</b>	35.836244	1.000000	53.000000	33768.420000	0.000000	2.000000	0.000000
<b>75%</b>	61.161020	3.000000	71.000000	54296.402500	0.000000	2.000000	1.000000
<b>max</b>	71.981490	10.000000	89.000000	207249.100000	1.000000	3.000000	1.000000

In [20]: `# now that all the modifications have been made, export the prepared dataset`  
`df.to_csv(r'C:\Users\fahim\Documents\0_WGUDocuments\d208\1medical_clean-PREPAREDTASK1.csv')`

In [21]: `# identify the columns for numerical data`  
`NumericalData = df.select_dtypes(include = "number").columns`  
`print (NumericalData)`

```
Index(['Initial_days', 'Children', 'Age', 'Income', 'Marital', 'Gender', 'ReAdmis', 'VitD_levels', 'Doc_visits', 'Full_meals_eaten', 'vitD_supp', 'Soft_drink', 'Initial_admin', 'HighBlood', 'Stroke', 'Complication_risk', 'Overweight', 'Arthritis', 'Diabetes', 'Hyperlipidemia', 'BackPain', 'Anxiety', 'Allergic_rhinitis', 'Reflux_esophagitis', 'Asthma', 'Services', 'TotalCharge', 'Additional_charges', 'Timely_admis', 'Timely_treat', 'Timely_visits', 'Reliability', 'Options', 'Hrs_treat', 'Courteous', 'Active_listen'],
      dtype='object')
```

In [22]: `# create histogram plots of the identified numerica data`  
`fig = plt.figure(figsize=(10, 20))`  
`ax = df[NumericalData].hist(bins = 15, figsize=(15,15))`  
`plt.title('Numeric Data')`  
`fig.tight_layout(h_pad=5, w_pad=5)`  
`plt.show()`


<Figure size 720x1440 with 0 Axes>




```
In [23]: # create boxplots for the continuous variables: Initial_days, TotalCharge, and Doc_vis
sns.boxplot('Initial_days', data = df)
plt.show()
sns.boxplot('TotalCharge', data = df)
plt.show()
sns.boxplot('Doc_visits', data = df)
plt.show()
```

C:\Users\fahim\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.


```
warnings.warn(
```





```
C:\Users\fahim\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn\_dec  
orators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From v  
ersion 0.12, the only valid positional argument will be `data`, and passing other arg  
uments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(
```




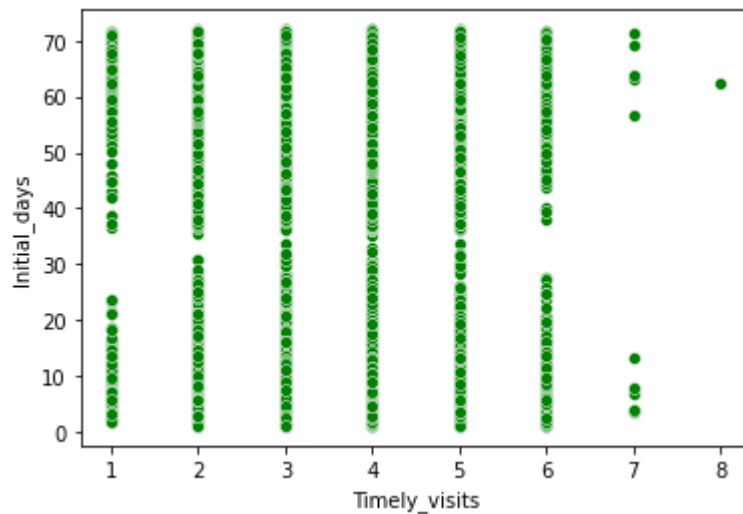
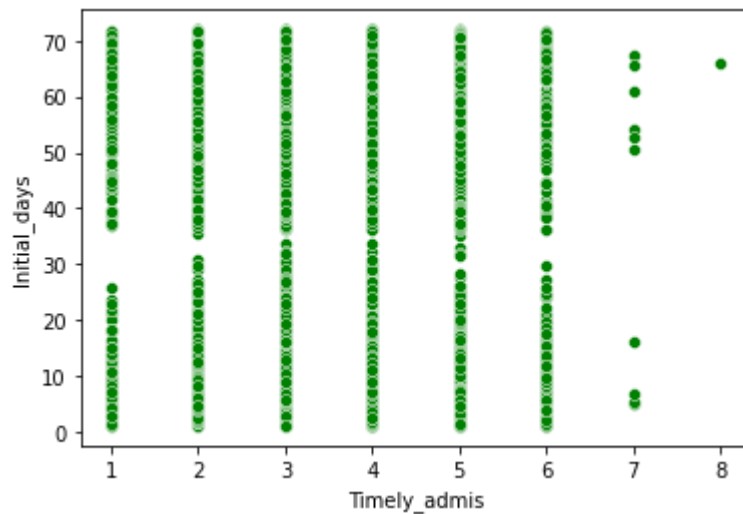

```
C:\Users\fahim\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn\_dec  
orators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From v  
ersion 0.12, the only valid positional argument will be `data`, and passing other arg  
uments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(
```



```
In [24]: # Create scatterplots to demonstrate and identify relationships between the target var
sns.scatterplot(x=df['Doc_visits'],y=df['Initial_days'],color='green')
plt.show();
sns.scatterplot(x=df['Initial_admin'],y=df['Initial_days'],color='green')
plt.show();
sns.scatterplot(x=df['ReAdmis'],y=df['Initial_days'],color='green')
plt.show();
sns.scatterplot(x=df['Complication_risk'],y=df['Initial_days'],color='green')
plt.show();
sns.scatterplot(x=df['Stroke'],y=df['Initial_days'],color='green')
plt.show();
sns.scatterplot(x=df['Diabetes'],y=df['Initial_days'],color='green')
plt.show();
sns.scatterplot(x=df['Timely_admis'],y=df['Initial_days'],color='green')
plt.show();
sns.scatterplot(x=df['Timely_visits'],y=df['Initial_days'],color='green')
plt.show();
df['intercept'] = 1
lm_initialdays = sm.OLS(df['Initial_days'],df[['Age', 'ReAdmis', 'Doc_visits', 'Initial_admin']])
```







C:\Users\fahim\AppData\Local\Temp\ipykernel\_21320\3725291535.py:18: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

`df['intercept'] = 1`

In [25]: `# retrieve the OLS Regression Results`

```
print(lm_initialdays.summary())
```

## OLS Regression Results

=====						
Dep. Variable:	Initial_days	R-squared:	0.994			
Model:	OLS	Adj. R-squared:	0.994			
Method:	Least Squares	F-statistic:	9.209e+04			
Date:	Mon, 10 Oct 2022	Prob (F-statistic):	0.00			
Time:	17:21:13	Log-Likelihood:	-21027.			
No. Observations:	10000	AIC:	4.209e+04			
Df Residuals:	9980	BIC:	4.224e+04			
Df Model:	19					
Covariance Type:	nonrobust					
=====						
=						
	coef	std err	t	P> t	[0.025	0.97
-----						
5]						
-----						
-						
Age	0.0197	0.003	6.780	0.000	0.014	0.02
5						
ReAdmis	0.9935	0.078	12.755	0.000	0.841	1.14
6						
Doc_visits	0.0197	0.019	1.037	0.300	-0.018	0.05
7						
Initial_admin	-3.3450	0.024	-137.279	0.000	-3.393	-3.29
7						
HighBlood	-0.4696	0.112	-4.175	0.000	-0.690	-0.24
9						
Stroke	0.0806	0.050	1.617	0.106	-0.017	0.17
8						
Complication_risk	-2.7188	0.028	-98.534	0.000	-2.773	-2.66
5						
Diabetes	-0.9084	0.045	-20.395	0.000	-0.996	-0.82
1						
Anxiety	-1.0138	0.043	-23.827	0.000	-1.097	-0.93
0						
Allergic_rhinitis	-0.8316	0.041	-20.459	0.000	-0.911	-0.75
2						
Reflux_esophagitis	-0.7705	0.040	-19.103	0.000	-0.850	-0.69
1						
Asthma	0.0236	0.044	0.540	0.589	-0.062	0.10
9						
Services	0.0096	0.020	0.475	0.635	-0.030	0.04
9						
TotalCharge	0.0119	1.73e-05	687.854	0.000	0.012	0.01
2						
Additional_charges	-9.473e-05	1.22e-05	-7.786	0.000	-0.000	-7.09e-0
5						
Timely_admis	-0.0339	0.027	-1.279	0.201	-0.086	0.01
8						
Timely_treat	0.0089	0.026	0.347	0.729	-0.041	0.05
9						
Hrs_treat	0.0138	0.022	0.629	0.529	-0.029	0.05
7						
Active_listen	0.0236	0.020	1.158	0.247	-0.016	0.06
4						
intercept	-14.5799	0.186	-78.567	0.000	-14.944	-14.21
6						
=====						
Omnibus:	284.824	Durbin-Watson:	2.008			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	223.140			



Skew: 0.283 Prob(JB): 3.51e-49  
 Kurtosis: 2.535 Cond. No. 1.47e+05  
 =====

## Notes:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.47e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [26]: # to address the strong multicollinearity, create heatmap and correlation matrix
medical_heatmap = df[['Initial_days', 'Age', 'ReAdmis', 'Doc_visits', 'Initial_admin',
#Initial model heatmap
sns.heatmap(medical_heatmap.corr(), annot=False)
plt.show
medical_heatmap.corr()
```

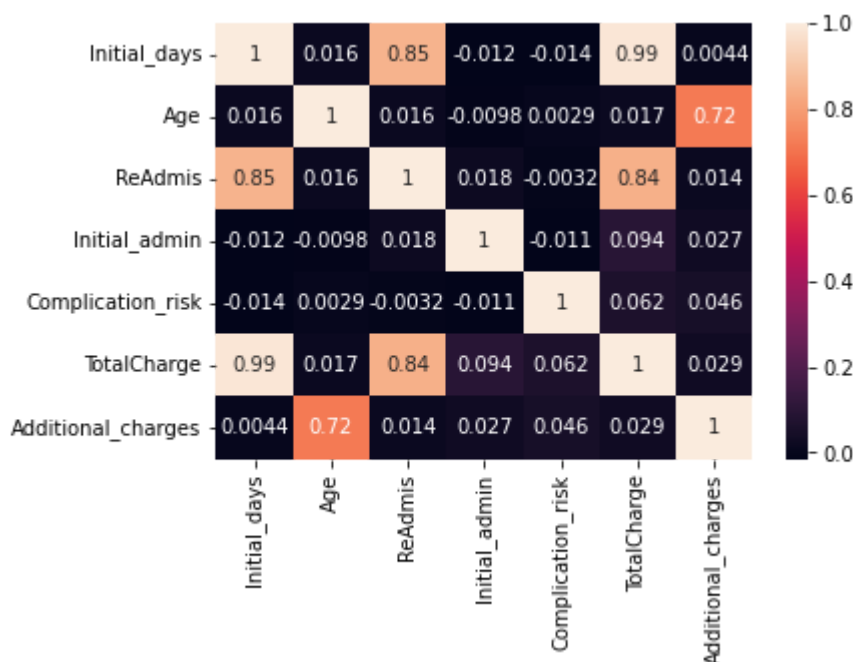
Out[26]:

	Initial_days	Age	ReAdmis	Doc_visits	Initial_admin	HighBlood	Stroke
Initial_days	1.000000	0.016264	0.850862	-0.006754	-0.012058	-0.006333	-0.002043
Age	0.016264	1.000000	0.015810	0.006898	-0.009763	0.007147	0.012035
ReAdmis	0.850862	0.015810	1.000000	0.000246	0.017522	0.002270	0.000918
Doc_visits	-0.006754	0.006898	0.000246	1.000000	0.012518	0.008967	-0.002230
Initial_admin	-0.012058	-0.009763	0.017522	0.012518	1.000000	0.001369	-0.008856
HighBlood	-0.006333	0.007147	0.002270	0.008967	0.001369	1.000000	0.007568
Stroke	-0.002043	0.012035	0.000918	-0.002230	-0.008856	0.007568	1.000000
Complication_risk	-0.014294	0.002887	-0.003236	0.012306	-0.011229	0.021368	0.001119
Diabetes	-0.002411	0.003694	-0.003058	0.012781	-0.009667	-0.005858	0.005792
Anxiety	0.011908	0.006130	0.002406	-0.001684	0.008305	0.008303	-0.013801
Allergic_rhinitis	0.003635	0.012092	-0.004651	0.002920	-0.005741	0.011709	-0.004837
Reflux_esophagitis	0.012237	-0.019609	0.005422	-0.005330	-0.004618	0.001150	-0.000054
Asthma	-0.013496	0.009229	-0.017133	-0.017989	-0.005956	0.006174	0.002443
Services	-0.007448	0.012016	-0.005578	-0.010785	0.003836	-0.003016	-0.016236
TotalCharge	0.987640	0.016876	0.843726	-0.005043	0.094157	0.019910	-0.003694
Additional_charges	0.004409	0.716854	0.013620	0.008072	0.026720	0.654316	0.035140
Timely_admis	-0.022258	0.005552	-0.016785	0.003680	0.006172	-0.011017	0.001948
Timely_treat	-0.007738	0.003967	-0.002423	0.006024	0.011959	-0.007745	-0.007706
Hrs_treat	-0.011752	-0.002087	-0.016894	0.012530	0.016487	-0.002369	0.004282
Active_listen	-0.008034	-0.003367	-0.016740	0.004571	-0.003092	0.002601	0.000040



```
In [27]: # to narrow the results, remove the diagnosis and survey variables and create a reduced
medical_heatmap = df[['Initial_days', 'Age', 'ReAdmis', 'Initial_admin', 'Complication_risk', 'TotalCharge', 'Additional_charges']]
sns.heatmap(medical_heatmap.corr(), annot=True)
plt.show
```

```
Out[27]: <function matplotlib.pyplot.show(close=None, block=None)>
```



```
In [28]: # create the reduced multiple regression model
df['intercept'] = 1
lm_initialdays_reduced = sm.OLS(df['Initial_days'], df[['Age', 'ReAdmis', 'Initial_admin', 'Complication_risk', 'TotalCharge', 'Additional_charges']])
print(lm_initialdays_reduced.summary())
```

## OLS Regression Results

```

=====
Dep. Variable:          Initial_days    R-squared:                0.993
Model:                  OLS             Adj. R-squared:           0.993
Method:                 Least Squares    F-statistic:             2.478e+05
Date:                   Mon, 10 Oct 2022  Prob (F-statistic):       0.00
Time:                   17:45:37         Log-Likelihood:          -21843.
No. Observations:       10000           AIC:                    4.370e+04
Df Residuals:           9993           BIC:                    4.375e+04
Df Model:                6
Covariance Type:        nonrobust
=====

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.97
5]
-----
-
Age              0.0308      0.001     20.537      0.000      0.028      0.03
4
ReAdmis          1.2333      0.084     14.650      0.000      1.068      1.39
8
Initial_admin    -3.3141      0.026    -126.496      0.000     -3.365     -3.26
3
Complication_risk -2.6898      0.030     -90.460      0.000     -2.748     -2.63
2
TotalCharge       0.0119    1.87e-05     633.739      0.000      0.012      0.01
2
Additional_charges -0.0001    4.73e-06     -30.383      0.000     -0.000     -0.00
0
intercept        -15.6501      0.120    -130.377      0.000     -15.885     -15.41
5
=====
Omnibus:             138.342    Durbin-Watson:           1.991
Prob(Omnibus):        0.000    Jarque-Bera (JB):        112.748
Skew:                 0.189    Prob(JB):                3.29e-25
Kurtosis:             2.642    Cond. No.:               8.98e+04
=====

```

## Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 8.98e+04. This might indicate that there are strong multicollinearity or other numerical problems.

C:\Users\fahim\AppData\Local\Temp\ipykernel\_21320\1666496053.py:2: SettingWithCopyWarning:


A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
df['intercept'] = 1

```
In [29]: # Load cleaned data to use for the residual plot; we will name the dataframe "regression_df"
regression_df = pd.read_csv(r'C:\Users\fahim\Documents\0_WGUDocuments\d208\1medical_cl
```

```
In [31]: # create the residual plot
regression_df['intercept'] = 1
residuals = regression_df['Initial_days'] - lm_initialdays_reduced.predict(regression_
```

```
sns.scatterplot(x=regression_df['TotalCharge'],y=residuals,color='green')  
plt.show();
```



In [ ]: