

```
In [1]: # import packages that will be used for the logistics regression analysis
import pylab
import seaborn as sb
sb.set(style="white")
sb.set(style="whitegrid", color_codes=True)
import sklearn
from sklearn.metrics import confusion_matrix
from sklearn import preprocessing
from sklearn.decomposition import PCA
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.metrics import classification_report
from sklearn import metrics
import matplotlib.pyplot as plt
plt.rc("font", size=14)
import numpy as np
import scipy.stats as stats
import statsmodels.api as sm
import statsmodels.formula.api as smf
from IPython.core.display import HTML
from IPython.display import display
import pandas as pd
from pandas import Series, DataFrame
from sklearn.metrics import classification_report, confusion_matrix
from imblearn.over_sampling import SMOTE

# import data set that will be used for the logistics regression analysis
pd.set_option('display.max_columns', None)
df = pd.read_csv (r'C:\Users\fahim\Documents\0_WGUDocuments\d208\1medical_clean.csv')

# rename the item columns accordingly
df.rename(columns={'Item1':'Timely_admis','Item2':'Timely_treat',
'Item3':'Timely_visits','Item4':'Reliability',
'Item5':'Options','Item6':'Hrs_treat',
'Item7':'Courteous','Item8':'Active_listen'},inplace=True)
df.head()
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 50 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CaseOrder              10000 non-null  int64
1   Customer_id            10000 non-null  object
2   Interaction             10000 non-null  object
3   UID                    10000 non-null  object
4   City                   10000 non-null  object
5   State                  10000 non-null  object
6   County                 10000 non-null  object
7   Zip                    10000 non-null  int64
8   Lat                    10000 non-null  float64
9   Lng                    10000 non-null  float64
10  Population              10000 non-null  int64
11  Area                    10000 non-null  object
12  TimeZone                10000 non-null  object
13  Job                     10000 non-null  object
14  Children                10000 non-null  int64
15  Age                     10000 non-null  int64
16  Income                  10000 non-null  float64
17  Marital                 10000 non-null  object
18  Gender                  10000 non-null  object
19  ReAdmis                 10000 non-null  object
20  VitD_levels             10000 non-null  float64
21  Doc_visits              10000 non-null  int64
22  Full_meals_eaten        10000 non-null  int64
23  vitD_supp               10000 non-null  int64
24  Soft_drink              10000 non-null  object
25  Initial_admin           10000 non-null  object
26  HighBlood               10000 non-null  object
27  Stroke                  10000 non-null  object
28  Complication_risk       10000 non-null  object
29  Overweight              10000 non-null  object
30  Arthritis               10000 non-null  object
31  Diabetes                10000 non-null  object
32  Hyperlipidemia          10000 non-null  object
33  BackPain                10000 non-null  object
34  Anxiety                 10000 non-null  object
35  Allergic_rhinitis       10000 non-null  object
36  Reflux_esophagitis      10000 non-null  object
37  Asthma                  10000 non-null  object
38  Services                 10000 non-null  object
39  Initial_days            10000 non-null  float64

```

```
40 TotalCharge      10000 non-null float64
41 Additional_charges 10000 non-null float64
42 Timely_admis     10000 non-null int64
43 Timely_treat     10000 non-null int64
44 Timely_visits    10000 non-null int64
45 Reliability      10000 non-null int64
46 Options          10000 non-null int64
47 Hrs_treat        10000 non-null int64
48 Courteous        10000 non-null int64
49 Active_listen    10000 non-null int64
```

```
dtypes: float64(7), int64(16), object(27)
```

```
memory usage: 3.8+ MB
```

```
In [2]: # drop all the demographic columns we don't need for this logistics regression analysis
df.drop(['City', 'State', 'County', 'Area', 'Zip', 'Lat', 'Lng', 'Population', 'TimeZone', 'Additional_charges', 'TotalCharge', 'I
# verify that all the columns were dropped before proceeding
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 34 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Children              10000 non-null  int64
1   Age                   10000 non-null  int64
2   Income                10000 non-null  float64
3   Marital               10000 non-null  object
4   Gender                10000 non-null  object
5   ReAdmis               10000 non-null  object
6   VitD_levels           10000 non-null  float64
7   Doc_visits            10000 non-null  int64
8   Full_meals_eaten      10000 non-null  int64
9   vitD_supp             10000 non-null  int64
10  Soft_drink             10000 non-null  object
11  Initial_admin          10000 non-null  object
12  HighBlood              10000 non-null  object
13  Stroke                 10000 non-null  object
14  Complication_risk      10000 non-null  object
15  Overweight             10000 non-null  object
16  Arthritis              10000 non-null  object
17  Diabetes               10000 non-null  object
18  Hyperlipidemia         10000 non-null  object
19  BackPain               10000 non-null  object
20  Anxiety                10000 non-null  object
21  Allergic_rhinitis      10000 non-null  object
22  Reflux_esophagitis     10000 non-null  object
23  Asthma                 10000 non-null  object
24  Services               10000 non-null  object
25  Initial_days           10000 non-null  float64
26  Timely_admis           10000 non-null  int64
27  Timely_treat           10000 non-null  int64
28  Timely_visits          10000 non-null  int64
29  Reliability            10000 non-null  int64
30  Options                10000 non-null  int64
31  Hrs_treat              10000 non-null  int64
32  Courteous              10000 non-null  int64
33  Active_listen          10000 non-null  int64
dtypes: float64(3), int64(13), object(18)
memory usage: 2.6+ MB

```

```

In [3]: #check if there is any duplicate data entries present in columns
df[df.duplicated()]

```

Out[3]: **Children Age Income Marital Gender ReAdmis VitD_levels Doc_visits Full_meals_eaten vitD_supp Soft_drink Initial_admin HighBlood**

In [4]: *# check if there are any duplicated columns in the data set - if there are none then the output should be False*
`df.columns.duplicated().any()`

Out[4]: False

In [5]: *# check if there are any duplicated rows in the data set - if there are none then the output should be False*
`df.duplicated().any()`

Out[5]: False

In [6]: *# convert categorical yes/no values to numeric 1/0 values*
`df = df.replace(to_replace = ['Yes', 'No'], value = [1, 0])`
`df`

Out[6]:

	Children	Age	Income	Marital	Gender	ReAdmis	VitD_levels	Doc_visits	Full_meals_eaten	vitD_supp	Soft_drink	Initial_admin	H
0	1	53	86575.93	Divorced	Male	0	19.141466	6	0	0	0	Emergency Admission	
1	3	51	46805.99	Married	Female	0	18.940352	4	2	1	0	Emergency Admission	
2	3	53	14370.14	Widowed	Female	0	18.057507	4	1	0	0	Elective Admission	
3	0	78	39741.49	Married	Male	0	16.576858	4	1	0	0	Elective Admission	
4	1	22	1209.56	Widowed	Female	0	17.439069	5	0	2	1	Elective Admission	
...
9995	2	25	45967.61	Widowed	Male	0	16.980860	4	2	1	0	Emergency Admission	
9996	4	87	14983.02	Widowed	Male	1	18.177020	5	0	0	0	Elective Admission	
9997	3	45	65917.81	Separated	Female	1	17.129070	4	2	0	1	Elective Admission	
9998	3	43	29702.32	Divorced	Male	1	19.910430	5	2	1	0	Emergency Admission	
9999	8	70	62682.63	Separated	Female	1	18.388620	5	0	1	0	Observation Admission	

10000 rows × 34 columns

```

In [7]: # convert the non-married Marital status values to "Married/Not Married", then convert "Married/Not Married" to "1/0"
#this will make the Marital variable easier to work with during regression analysis
df['Marital'] = df['Marital'].replace(['Divorced','Widowed','Separated','Never Married'],'Not Married')
df['Marital'] = df['Marital'].replace(['Married','Not Married'],[1,0])
df

```

Out[7]:

	Children	Age	Income	Marital	Gender	ReAdmis	VitD_levels	Doc_visits	Full_meals_eaten	vitD_supp	Soft_drink	Initial_admin	Higl
0	1	53	86575.93	0	Male	0	19.141466	6	0	0	0	Emergency Admission	
1	3	51	46805.99	1	Female	0	18.940352	4	2	1	0	Emergency Admission	
2	3	53	14370.14	0	Female	0	18.057507	4	1	0	0	Elective Admission	
3	0	78	39741.49	1	Male	0	16.576858	4	1	0	0	Elective Admission	
4	1	22	1209.56	0	Female	0	17.439069	5	0	2	1	Elective Admission	
...
9995	2	25	45967.61	0	Male	0	16.980860	4	2	1	0	Emergency Admission	
9996	4	87	14983.02	0	Male	1	18.177020	5	0	0	0	Elective Admission	
9997	3	45	65917.81	0	Female	1	17.129070	4	2	0	1	Elective Admission	
9998	3	43	29702.32	0	Male	1	19.910430	5	2	1	0	Emergency Admission	
9999	8	70	62682.63	0	Female	1	18.388620	5	0	1	0	Observation Admission	

10000 rows × 34 columns

```
In [8]: # Showcase the unique values for the Services variable
df['Gender'].unique()
```

```
Out[8]: array(['Male', 'Female', 'Nonbinary'], dtype=object)
```

```
In [9]: #convert the non-Female gender values to "Female/non-female", then convert "Female/non-female" to "1/0"
df['Gender'] = df['Gender'].replace(['Male', 'Nonbinary'], 'non-female')
df['Gender'] = df['Gender'].replace(['Female', 'non-female'], [1,0])
df
```

Out[9]:

	Children	Age	Income	Marital	Gender	ReAdmis	VitD_levels	Doc_visits	Full_meals_eaten	vitD_supp	Soft_drink	Initial_admin	Higl
0	1	53	86575.93	0	0	0	19.141466	6	0	0	0	Emergency Admission	
1	3	51	46805.99	1	1	0	18.940352	4	2	1	0	Emergency Admission	
2	3	53	14370.14	0	1	0	18.057507	4	1	0	0	Elective Admission	
3	0	78	39741.49	1	0	0	16.576858	4	1	0	0	Elective Admission	
4	1	22	1209.56	0	1	0	17.439069	5	0	2	1	Elective Admission	
...
9995	2	25	45967.61	0	0	0	16.980860	4	2	1	0	Emergency Admission	
9996	4	87	14983.02	0	0	1	18.177020	5	0	0	0	Elective Admission	
9997	3	45	65917.81	0	1	1	17.129070	4	2	0	1	Elective Admission	
9998	3	43	29702.32	0	0	1	19.910430	5	2	1	0	Emergency Admission	
9999	8	70	62682.63	0	1	1	18.388620	5	0	1	0	Observation Admission	

10000 rows × 34 columns

```
In [10]: # Showcase the unique values for the Services variable
df['Services'].unique()
```

```
Out[10]: array(['Blood Work', 'Intravenous', 'CT Scan', 'MRI'], dtype=object)
```

```
In [11]: # Drop the services variable since these values cannot be condensed
df.drop(['Services'],axis = 1,inplace=True)
```



```
In [12]: # Showcase the unique values for the Complication_risk variable  
df['Complication_risk'].unique()
```

```
Out[12]: array(['Medium', 'High', 'Low'], dtype=object)
```

```
In [13]: # Drop the services variable since these values cannot be condensed  
df.drop(['Complication_risk'],axis = 1,inplace=True)
```

```
In [14]: # Showcase the unique values for the Initial_admin variable  
df['Initial_admin'].unique()
```

```
Out[14]: array(['Emergency Admission', 'Elective Admission',  
               'Observation Admission'], dtype=object)
```

```
In [15]: # convert the non-emergency admission status values to "Emergency Admission/non-Emergency Admission", then convert "Eme  
#this will make the Marital variable easier to work with during regression analysis  
df['Initial_admin'] = df['Initial_admin'].replace(['Elective Admission','Observation Admission'],'non-Emergency Admissi  
df['Initial_admin'] = df['Initial_admin'].replace(['Emergency Admission','non-Emergency Admission'],[1,0])  
df
```

Out[15]:

	Children	Age	Income	Marital	Gender	ReAdmis	VitD_levels	Doc_visits	Full_meals_eaten	vitD_supp	Soft_drink	Initial_admin	Higl
0	1	53	86575.93	0	0	0	19.141466	6	0	0	0	1	
1	3	51	46805.99	1	1	0	18.940352	4	2	1	0	1	
2	3	53	14370.14	0	1	0	18.057507	4	1	0	0	0	
3	0	78	39741.49	1	0	0	16.576858	4	1	0	0	0	
4	1	22	1209.56	0	1	0	17.439069	5	0	2	1	0	
...
9995	2	25	45967.61	0	0	0	16.980860	4	2	1	0	1	
9996	4	87	14983.02	0	0	1	18.177020	5	0	0	0	0	
9997	3	45	65917.81	0	1	1	17.129070	4	2	0	1	0	
9998	3	43	29702.32	0	0	1	19.910430	5	2	1	0	1	
9999	8	70	62682.63	0	1	1	18.388620	5	0	1	0	0	

10000 rows × 32 columns

In [16]: *# describe the dataframe and showcase summary statistics of the variables*
 df.describe()

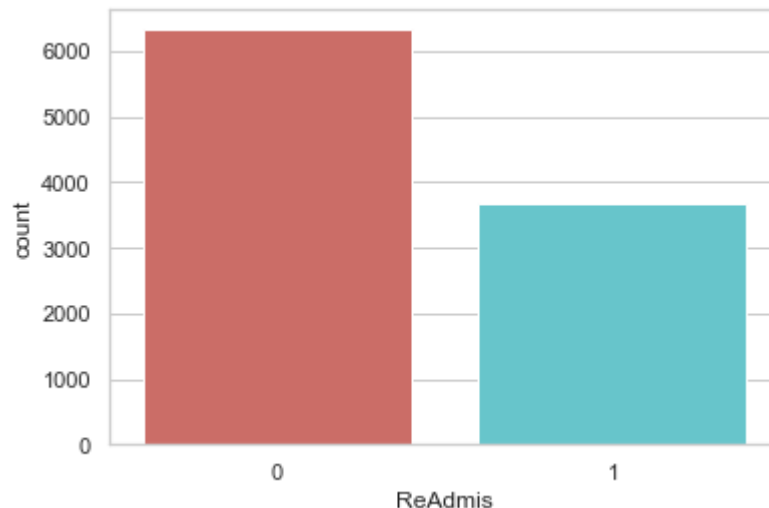
Out[16]:

	Children	Age	Income	Marital	Gender	ReAdmis	VitD_levels	Doc_visits	Full_meals_eaten
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	2.097200	53.511700	40490.495160	0.202300	0.501800	0.366900	17.964262	5.012200	1.001400
std	2.163659	20.638538	28521.153293	0.401735	0.500022	0.481983	2.017231	1.045734	1.008117
min	0.000000	18.000000	154.080000	0.000000	0.000000	0.000000	9.806483	1.000000	0.000000
25%	0.000000	36.000000	19598.775000	0.000000	0.000000	0.000000	16.626439	4.000000	0.000000
50%	1.000000	53.000000	33768.420000	0.000000	1.000000	0.000000	17.951122	5.000000	1.000000
75%	3.000000	71.000000	54296.402500	0.000000	1.000000	1.000000	19.347963	6.000000	2.000000
max	10.000000	89.000000	207249.100000	1.000000	1.000000	1.000000	26.394449	9.000000	7.000000

```
In [17]: # now that all the modifications have been made, export the prepared dataset
df.to_csv(r'C:\Users\fahim\Documents\0_WGUDocuments\d208\2medical_clean-PREPAREDTASK2_12-24-2022.csv')
```

```
In [18]: # Start to visualize the data, including univariate and bivariate analyses
# begin by visualizing the target variable, ReAdmis
print(df['ReAdmis'].value_counts())
sb.countplot(x='ReAdmis', data=df, palette='hls')
plt.show()
```

```
0    6331
1    3669
Name: ReAdmis, dtype: int64
```

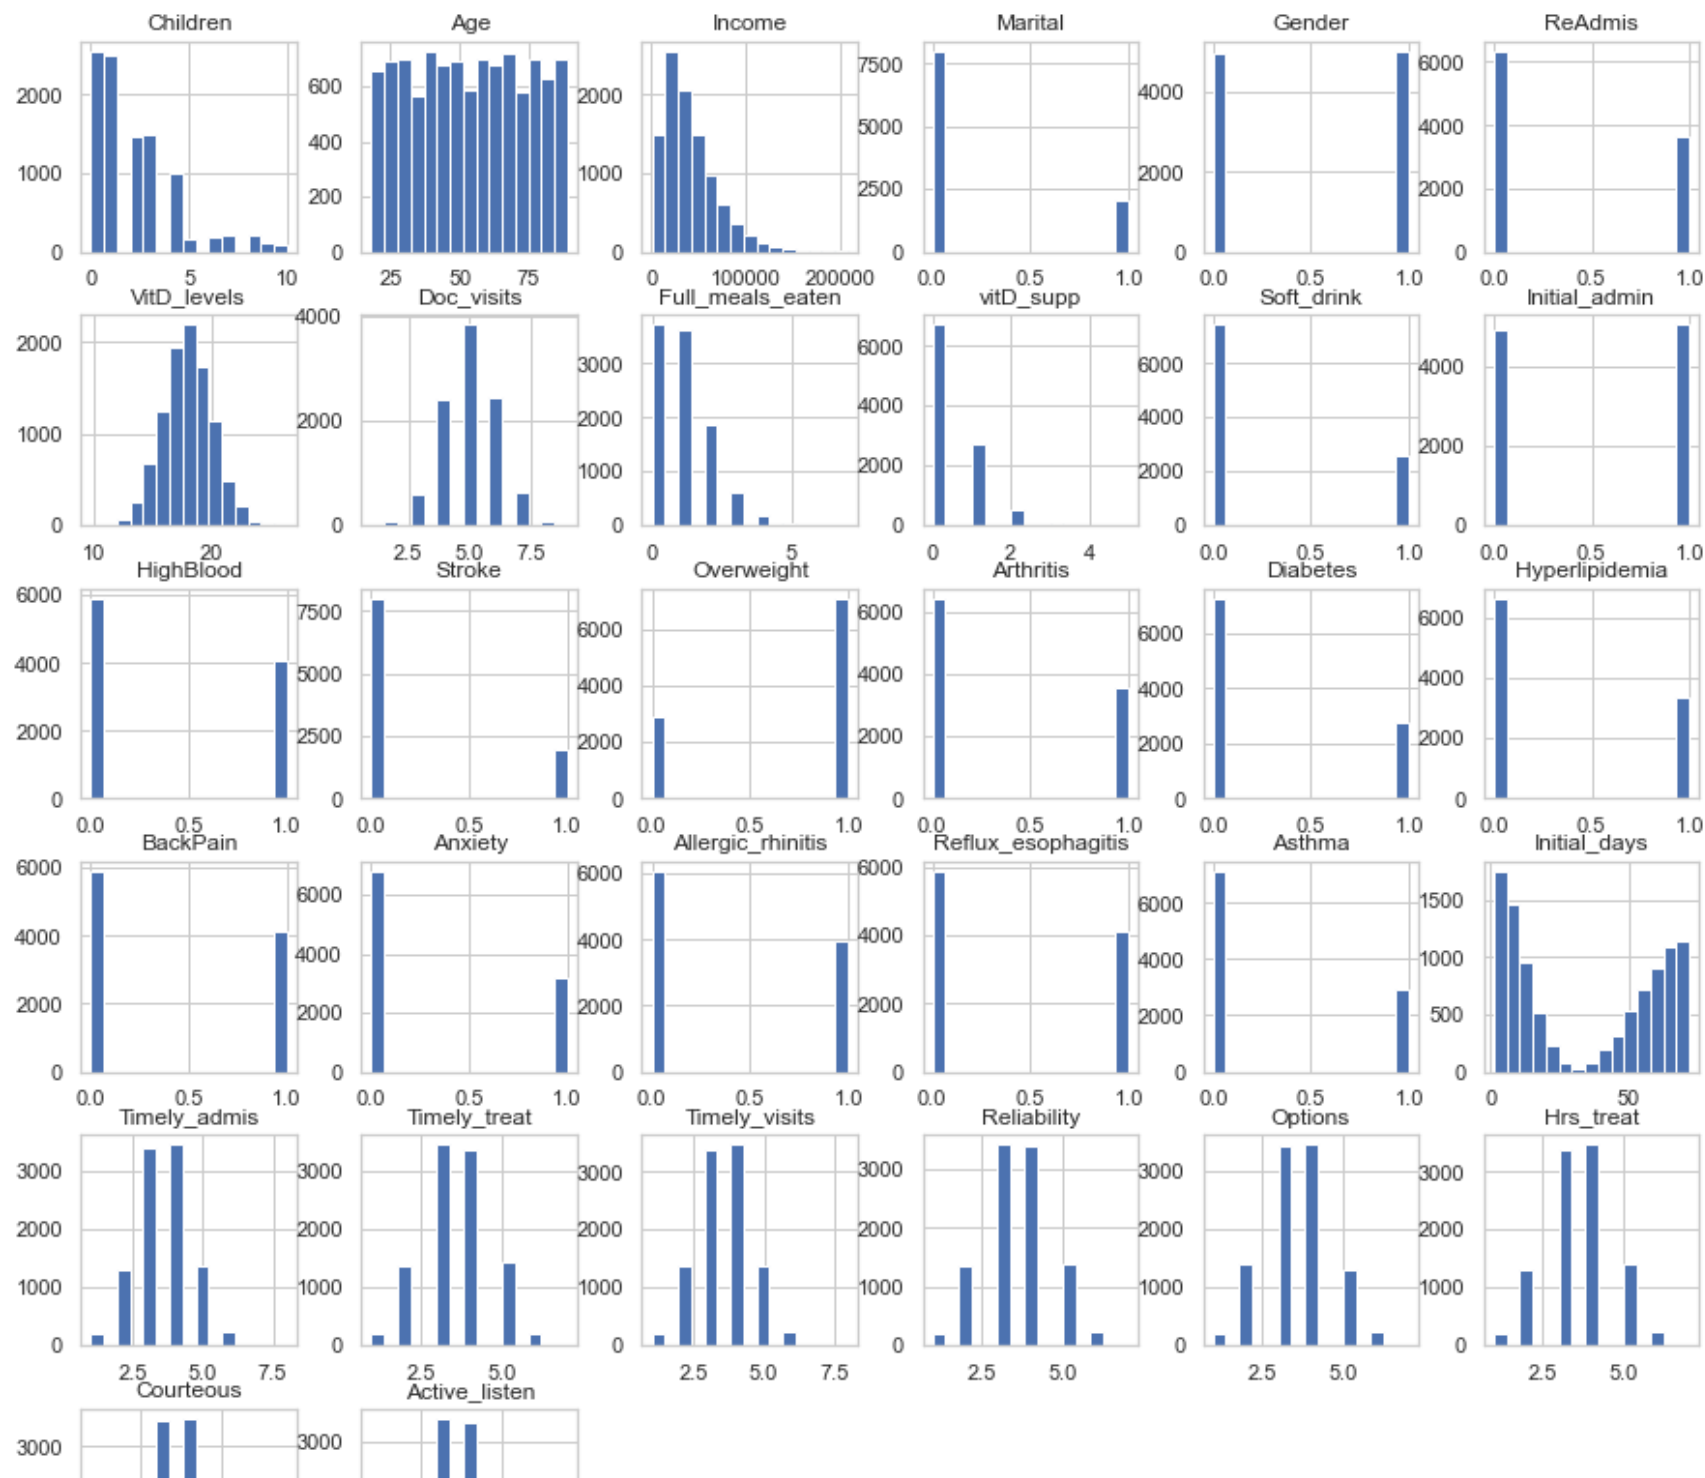



```
In [19]: # identify the columns for variables
Variables = df.select_dtypes(include = "number").columns
print (Variables)
```

```
Index(['Children', 'Age', 'Income', 'Marital', 'Gender', 'ReAdmis',
      'VitD_levels', 'Doc_visits', 'Full_meals_eaten', 'vitD_supp',
      'Soft_drink', 'Initial_admin', 'HighBlood', 'Stroke', 'Overweight',
      'Arthritis', 'Diabetes', 'Hyperlipidemia', 'BackPain', 'Anxiety',
      'Allergic_rhinitis', 'Reflux_esophagitis', 'Asthma', 'Initial_days',
      'Timely_admis', 'Timely_treat', 'Timely_visits', 'Reliability',
      'Options', 'Hrs_treat', 'Courteous', 'Active_listen'],
      dtype='object')
```

```
In [20]: # create histogram plots of the identified predictor variables
fig = plt.figure(figsize=(15, 15))
ax = df[Variables].hist(bins = 15, figsize=(15,15))
plt.title('Variables')
fig.tight_layout(h_pad=5, w_pad=5)
plt.show()
```

<Figure size 1080x1080 with 0 Axes>









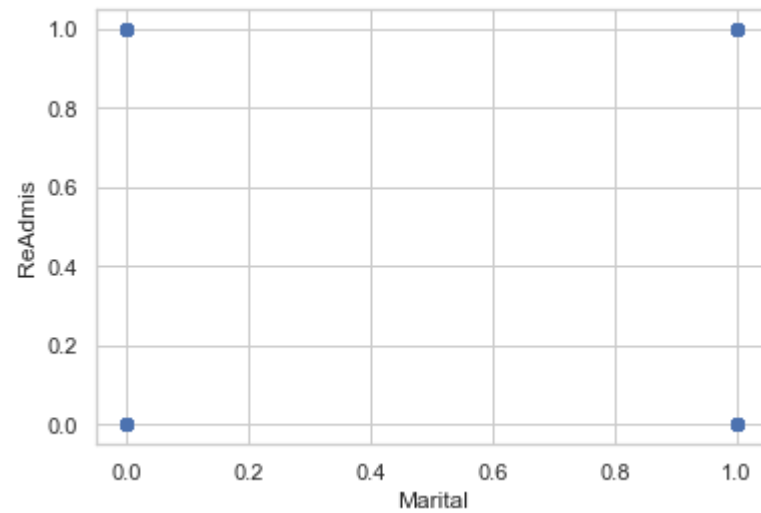
In [21]: *#selecting the target variable and showcasing the bivariate statistics*

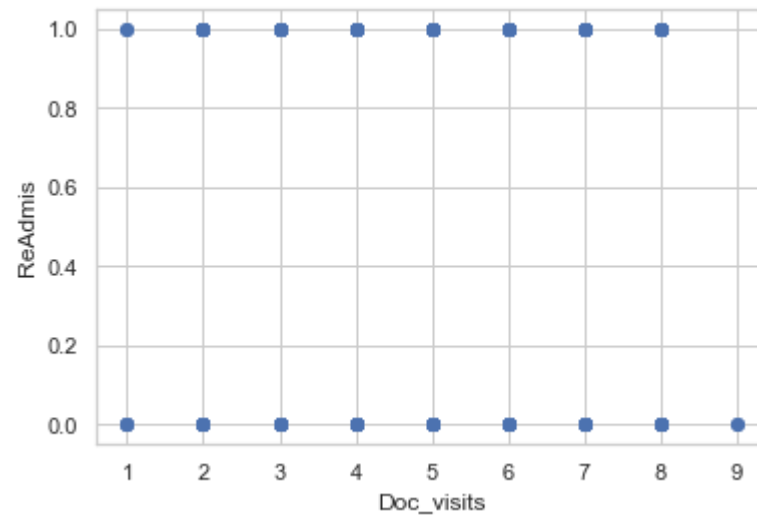
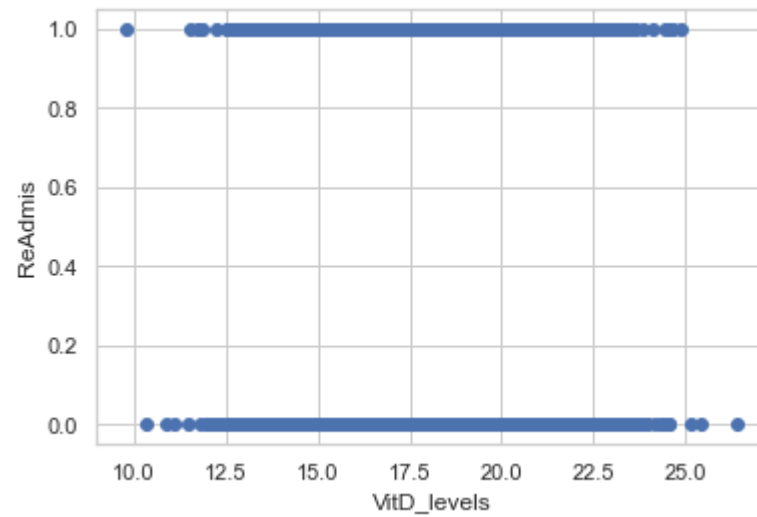
```
X=df.drop('ReAdmis',inplace=False, axis=1)
y=df['ReAdmis']


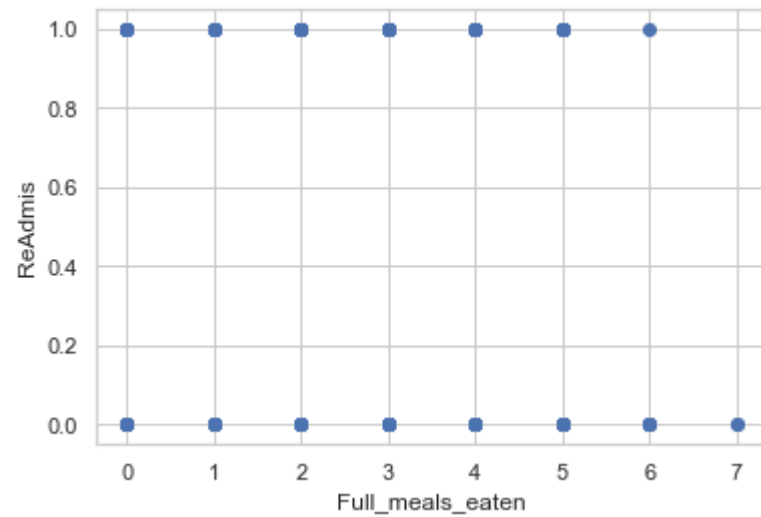
for column in X.columns:
    plt.scatter(X[column],y)
    plt.xlabel(column)
    plt.ylabel('ReAdmis')
    plt.show()
```





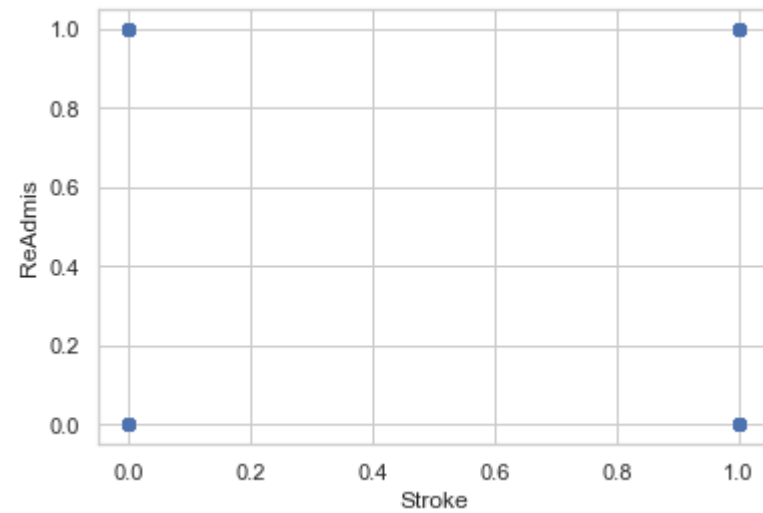







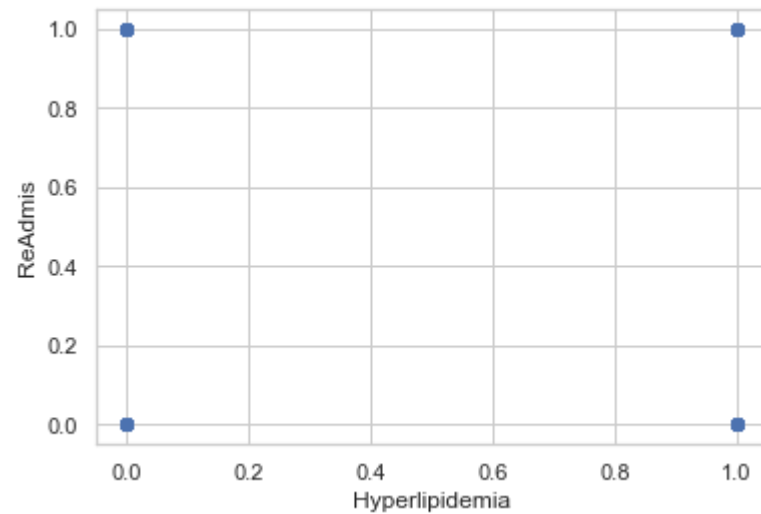




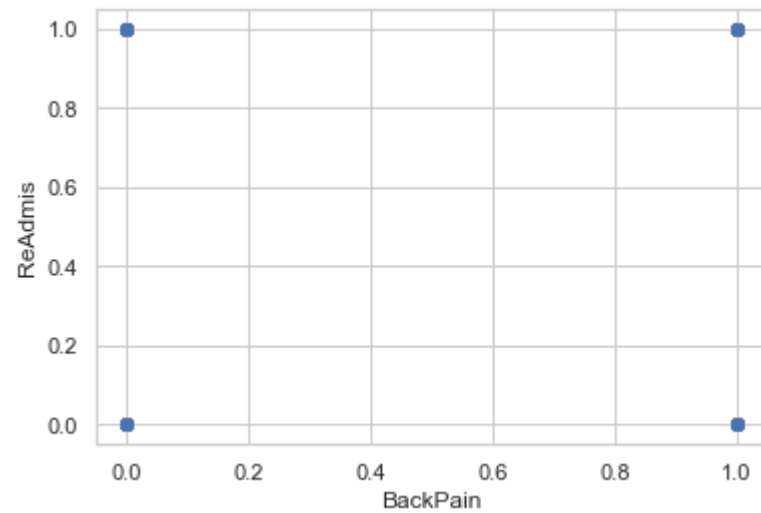


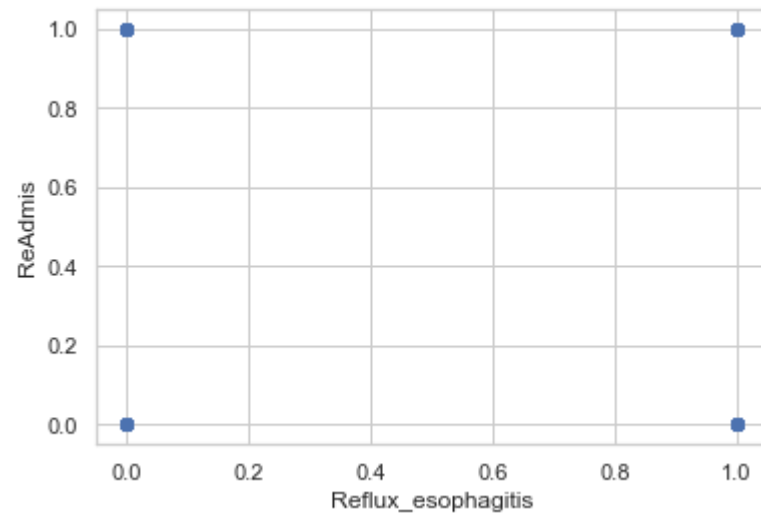
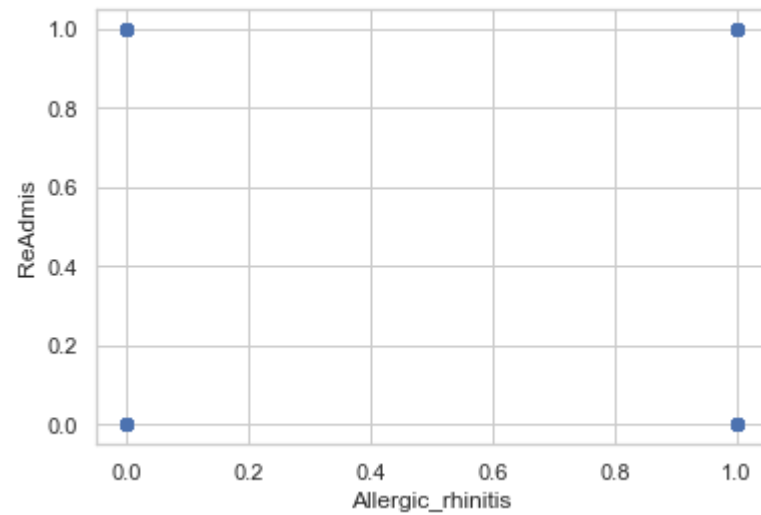















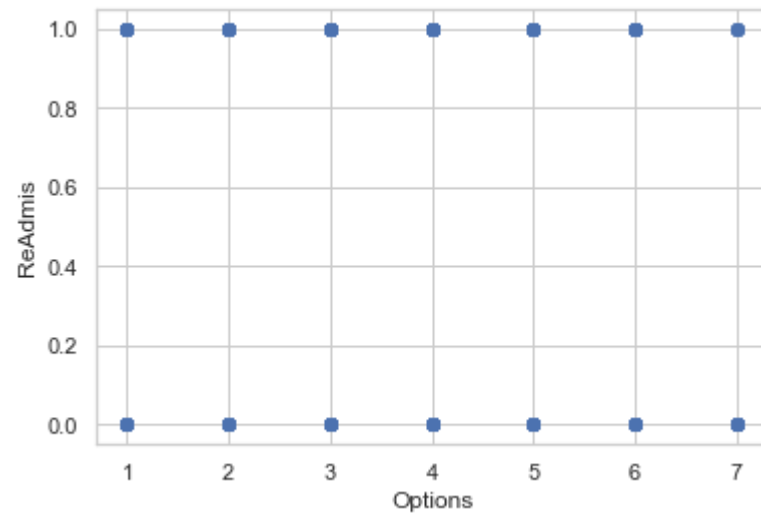
















```
In [22]: # create the initial logistics model
log_reg_results = sm.Logit(df["ReAdmis"], df[['Children', 'Age', 'Income', 'Marital', 'Gender', 'VitD_levels', 'Doc_vis
print(log_reg_results.summary())
```

Optimization terminated successfully.

Current function value: 0.107566

Iterations 11

Logit Regression Results

=====						
Dep. Variable:	ReAdmis	No. Observations:	10000			
Model:	Logit	Df Residuals:	9969			
Method:	MLE	Df Model:	30			
Date:	Sat, 24 Dec 2022	Pseudo R-squ.:	0.8363			
Time:	23:42:39	Log-Likelihood:	-1075.7			
converged:	True	LL-Null:	-6572.9			
Covariance Type:	nonrobust	LLR p-value:	0.000			
=====						
	coef	std err	z	P> z	[0.025	0.975]

Children	0.0050	0.025	0.200	0.841	-0.044	0.054
Age	-0.0139	0.003	-5.094	0.000	-0.019	-0.009
Income	-6.555e-06	1.92e-06	-3.422	0.001	-1.03e-05	-2.8e-06
Marital	-0.0696	0.141	-0.495	0.620	-0.345	0.206
Gender	-0.3728	0.112	-3.338	0.001	-0.592	-0.154
VitD_levels	-0.4593	0.026	-17.936	0.000	-0.509	-0.409
Doc_visits	-0.4609	0.053	-8.763	0.000	-0.564	-0.358
Full_meals_eaten	-0.0743	0.056	-1.338	0.181	-0.183	0.035
vitD_supp	-0.1947	0.086	-2.254	0.024	-0.364	-0.025
Soft_drink	-0.0081	0.128	-0.064	0.949	-0.258	0.242
Initial_admin	0.4679	0.111	4.197	0.000	0.249	0.686
HighBlood	0.1121	0.113	0.990	0.322	-0.110	0.334
Stroke	0.2222	0.140	1.584	0.113	-0.053	0.497
Overweight	-0.4160	0.122	-3.416	0.001	-0.655	-0.177
Arthritis	-0.4974	0.115	-4.320	0.000	-0.723	-0.272
Diabetes	-0.1160	0.123	-0.944	0.345	-0.357	0.125
Hyperlipidemia	-0.0781	0.118	-0.664	0.507	-0.308	0.152
BackPain	-0.0150	0.112	-0.134	0.894	-0.235	0.205
Anxiety	-0.3797	0.119	-3.198	0.001	-0.612	-0.147
Allergic_rhinitis	-0.3806	0.113	-3.358	0.001	-0.603	-0.158
Reflux_esophagitis	-0.3800	0.113	-3.348	0.001	-0.602	-0.158
Asthma	-0.4210	0.123	-3.432	0.001	-0.661	-0.181
Initial_days	0.4063	0.013	31.118	0.000	0.381	0.432
Timely_admis	0.1473	0.081	1.828	0.068	-0.011	0.305
Timely_treat	0.0719	0.074	0.976	0.329	-0.072	0.216
Timely_visits	-0.3015	0.068	-4.402	0.000	-0.436	-0.167
Reliability	-0.5018	0.060	-8.365	0.000	-0.619	-0.384
Options	-0.9190	0.065	-14.239	0.000	-1.045	-0.792
Hrs_treat	-0.3222	0.068	-4.747	0.000	-0.455	-0.189
Courteous	-0.3539	0.061	-5.823	0.000	-0.473	-0.235

Active_listen	-0.3633	0.058	-6.211	0.000	-0.478	-0.249
---------------	---------	-------	--------	-------	--------	--------

=====

Possibly complete quasi-separation: A fraction 0.51 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

```
In [23]: # create the correlation matrix
matrix_df = pd.read_csv(r'C:\Users\fahim\Documents\0_WGUDocuments\d208\2medical_clean-PREPAREDTASK2_12-24-2022.csv')

matrix_df = matrix_df[['Children', 'Age', 'Income', 'Marital', 'Gender', 'VitD_levels', 'Doc_visits', 'Full_meals_eaten']]

X = matrix_df.iloc[:, 1:-1].values
y = matrix_df.iloc[:, -1].values
```

```
In [24]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)
```

```
In [25]: from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression(random_state = 0)
classifier.fit(X_train, y_train)
```

C:\Users\fahim\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\linear_model_logistic.py:444: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

```
Out[25]: ▼ LogisticRegression
LogisticRegression(random_state=0)
```


```
In [26]: y_pred = classifier.predict(X_test)
```

```
In [27]: #now create the confusion matrix for the initial model
from sklearn.metrics import confusion_matrix
matrix = confusion_matrix(y_test, y_pred)
print(matrix)
```

```
[[1189  73]
 [  50 688]]
```

```
In [28]: y_predict_test = classifier.predict(X_test)
new_matrix = confusion_matrix(y_test, y_predict_test)
sb.heatmap(new_matrix, annot=True)
```

```
Out[28]: <AxesSubplot:>
```



```
In [29]: #retrieve the classification report for the initial model
from sklearn.metrics import classification_report
print(classification_report(y_test, y_predict_test))
```

	precision	recall	f1-score	support
0	0.96	0.94	0.95	1262
1	0.90	0.93	0.92	738
accuracy			0.94	2000
macro avg	0.93	0.94	0.93	2000
weighted avg	0.94	0.94	0.94	2000

```
In [30]: #Create the reduced model with the variables that had a P value below .05 statistical significance level
log_reg_results2 = sm.Logit(df["ReAdmis"], df[['Age', 'Income', 'Gender', 'VitD_levels', 'Doc_visits', 'vitD_supp', 'Ini
print(log_reg_results2.summary())
```


Optimization terminated successfully.
 Current function value: 0.107964
 Iterations 11

Logit Regression Results

```

=====
Dep. Variable:          ReAdmis    No. Observations:          10000
Model:                  Logit      Df Residuals:              9979
Method:                 MLE        Df Model:                  20
Date:                   Sat, 24 Dec 2022    Pseudo R-squ.:          0.8357
Time:                   23:42:40           Log-Likelihood:         -1079.6
converged:              True          LL-Null:                 -6572.9
Covariance Type:        nonrobust         LLR p-value:             0.000
=====

```

	coef	std err	z	P> z	[0.025	0.975]
Age	-0.0138	0.003	-5.099	0.000	-0.019	-0.009
Income	-6.401e-06	1.91e-06	-3.350	0.001	-1.01e-05	-2.66e-06
Gender	-0.3735	0.111	-3.356	0.001	-0.592	-0.155
VitD_levels	-0.4562	0.025	-18.004	0.000	-0.506	-0.406
Doc_visits	-0.4580	0.052	-8.803	0.000	-0.560	-0.356
vitD_supp	-0.1902	0.086	-2.206	0.027	-0.359	-0.021
Initial_admin	0.4636	0.111	4.176	0.000	0.246	0.681
Overweight	-0.4015	0.121	-3.315	0.001	-0.639	-0.164
Arthritis	-0.5081	0.115	-4.430	0.000	-0.733	-0.283
Anxiety	-0.3776	0.118	-3.199	0.001	-0.609	-0.146
Allergic_rhinitis	-0.3773	0.113	-3.346	0.001	-0.598	-0.156
Reflux_esophagitis	-0.3738	0.113	-3.309	0.001	-0.595	-0.152
Asthma	-0.4248	0.122	-3.472	0.001	-0.665	-0.185
Initial_days	0.4039	0.013	31.204	0.000	0.379	0.429
Timely_admis	0.1828	0.072	2.529	0.011	0.041	0.324
Timely_visits	-0.2852	0.066	-4.299	0.000	-0.415	-0.155
Reliability	-0.4999	0.060	-8.393	0.000	-0.617	-0.383
Options	-0.9164	0.064	-14.346	0.000	-1.042	-0.791
Hrs_treat	-0.3167	0.067	-4.711	0.000	-0.448	-0.185
Courteous	-0.3487	0.060	-5.781	0.000	-0.467	-0.230
Active_listen	-0.3621	0.058	-6.230	0.000	-0.476	-0.248

```

=====

```

Possibly complete quasi-separation: A fraction 0.50 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

```

In [31]: # create the correlation matrix for the reduced model
matrix_df = pd.read_csv(r'C:\Users\fahim\Documents\0_WGUDocuments\d208\2medical_clean-PREPAREDTASK2_12-24-2022.csv')

```

```
matrix_df = matrix_df[['Age', 'Income', 'Gender', 'VitD_levels', 'Doc_visits', 'vitD_supp', 'Initial_admin', 'Overweight', 'Depression']]

X = matrix_df.iloc[:, 1:-1].values
y = matrix_df.iloc[:, -1].values
```

```
In [32]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)
```

```
In [33]: from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression(random_state = 0)
classifier.fit(X_train, y_train)
```

C:\Users\fahim\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\linear_model_logistic.py:444: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

```
Out[33]: LogisticRegression
LogisticRegression(random_state=0)
```


```
In [34]: y_pred = classifier.predict(X_test)
```

```
In [35]: #now create the confusion matrix for the reduced model
from sklearn.metrics import confusion_matrix
matrix = confusion_matrix(y_test, y_pred)
print(matrix)
```

```
[[1193  69]
 [  37 701]]
```

```
In [36]: y_predict_test = classifier.predict(X_test)
new_matrix = confusion_matrix(y_test, y_predict_test)
sb.heatmap(new_matrix, annot=True)
```


```
Out[36]: <AxesSubplot:>
```



```
In [37]: #retrieve the classificaiton report for the reduced model
from sklearn.metrics import classification_report
print(classification_report(y_test, y_predict_test))
```

	precision	recall	f1-score	support
0	0.97	0.95	0.96	1262
1	0.91	0.95	0.93	738
accuracy			0.95	2000
macro avg	0.94	0.95	0.94	2000
weighted avg	0.95	0.95	0.95	2000

```
In [38]: # plot ROC Curve
logit_roc_auc = roc_auc_score(y_test, classifier.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, classifier.predict_proba(X_test)[: ,1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' %logit_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



```
In [39]: # create an equation of the regression
print('Logit: {:.2f}'.format(logit_roc_auc))
equation = log_reg_results2.summary().tables[1]
print('Estimate [{}] as L = {}'.format(log_reg_results2.summary().tables[0][1][1]))
equation = pd.DataFrame(equation)
for i in equation.itertuples():
    print(' {:.3f} x ( {} ) '.format(i[0],i[1]))
```

```
Logit: 0.95
Estimate [Logit] as L =
+0.000 x ( )
+1.000 x ( Age )
+2.000 x ( Income )
+3.000 x ( Gender )
+4.000 x ( VitD_levels )
+5.000 x ( Doc_visits )
+6.000 x ( vitD_supp )
+7.000 x ( Initial_admin )
+8.000 x ( Overweight )
+9.000 x ( Arthritis )
+10.000 x ( Anxiety )
+11.000 x ( Allergic_rhinitis )
+12.000 x ( Reflux_esophagitis )
+13.000 x ( Asthma )
+14.000 x ( Initial_days )
+15.000 x ( Timely_admis )
+16.000 x ( Timely_visits )
+17.000 x ( Reliability )
+18.000 x ( Options )
+19.000 x ( Hrs_treat )
+20.000 x ( Courteous )
+21.000 x ( Active_listen )
```

```
In [40]: print(equation)
```

	0	1	2	3	4	5 \
0		coef	std err	z	P> z	[0.025
1	Age	-0.0138	0.003	-5.099	0.000	-0.019
2	Income	-6.401e-06	1.91e-06	-3.350	0.001	-1.01e-05
3	Gender	-0.3735	0.111	-3.356	0.001	-0.592
4	VitD_levels	-0.4562	0.025	-18.004	0.000	-0.506
5	Doc_visits	-0.4580	0.052	-8.803	0.000	-0.560
6	vitD_supp	-0.1902	0.086	-2.206	0.027	-0.359
7	Initial_admin	0.4636	0.111	4.176	0.000	0.246
8	Overweight	-0.4015	0.121	-3.315	0.001	-0.639
9	Arthritis	-0.5081	0.115	-4.430	0.000	-0.733
10	Anxiety	-0.3776	0.118	-3.199	0.001	-0.609
11	Allergic_rhinitis	-0.3773	0.113	-3.346	0.001	-0.598
12	Reflux_esophagitis	-0.3738	0.113	-3.309	0.001	-0.595
13	Asthma	-0.4248	0.122	-3.472	0.001	-0.665
14	Initial_days	0.4039	0.013	31.204	0.000	0.379
15	Timely_admis	0.1828	0.072	2.529	0.011	0.041
16	Timely_visits	-0.2852	0.066	-4.299	0.000	-0.415
17	Reliability	-0.4999	0.060	-8.393	0.000	-0.617
18	Options	-0.9164	0.064	-14.346	0.000	-1.042
19	Hrs_treat	-0.3167	0.067	-4.711	0.000	-0.448
20	Courteous	-0.3487	0.060	-5.781	0.000	-0.467
21	Active_listen	-0.3621	0.058	-6.230	0.000	-0.476

	6
0	0.975]
1	-0.009
2	-2.66e-06
3	-0.155
4	-0.406
5	-0.356
6	-0.021
7	0.681
8	-0.164
9	-0.283
10	-0.146
11	-0.156
12	-0.152
13	-0.185
14	0.429
15	0.324
16	-0.155
17	-0.383
18	-0.791
19	-0.185

20	-0.230
21	-0.248

```
In [41]: updated_equation = equation.drop(0)
```

```
In [42]: print(updated_equation)
```

	0	1	2	3	4	5 \
1	Age	-0.0138	0.003	-5.099	0.000	-0.019
2	Income	-6.401e-06	1.91e-06	-3.350	0.001	-1.01e-05
3	Gender	-0.3735	0.111	-3.356	0.001	-0.592
4	VitD_levels	-0.4562	0.025	-18.004	0.000	-0.506
5	Doc_visits	-0.4580	0.052	-8.803	0.000	-0.560
6	vitD_supp	-0.1902	0.086	-2.206	0.027	-0.359
7	Initial_admin	0.4636	0.111	4.176	0.000	0.246
8	Overweight	-0.4015	0.121	-3.315	0.001	-0.639
9	Arthritis	-0.5081	0.115	-4.430	0.000	-0.733
10	Anxiety	-0.3776	0.118	-3.199	0.001	-0.609
11	Allergic_rhinitis	-0.3773	0.113	-3.346	0.001	-0.598
12	Reflux_esophagitis	-0.3738	0.113	-3.309	0.001	-0.595
13	Asthma	-0.4248	0.122	-3.472	0.001	-0.665
14	Initial_days	0.4039	0.013	31.204	0.000	0.379
15	Timely_admis	0.1828	0.072	2.529	0.011	0.041
16	Timely_visits	-0.2852	0.066	-4.299	0.000	-0.415
17	Reliability	-0.4999	0.060	-8.393	0.000	-0.617
18	Options	-0.9164	0.064	-14.346	0.000	-1.042
19	Hrs_treat	-0.3167	0.067	-4.711	0.000	-0.448
20	Courteous	-0.3487	0.060	-5.781	0.000	-0.467
21	Active_listen	-0.3621	0.058	-6.230	0.000	-0.476

	6
1	-0.009
2	-2.66e-06
3	-0.155
4	-0.406
5	-0.356
6	-0.021
7	0.681
8	-0.164
9	-0.283
10	-0.146
11	-0.156
12	-0.152
13	-0.185
14	0.429
15	0.324
16	-0.155
17	-0.383
18	-0.791
19	-0.185
20	-0.230
21	-0.248


```
In [43]: # create an equation of the logistics regression
print('Logit: {:.2f}'.format(logit_roc_auc))
print('Estimate [{}] as L = '.format(log_reg_results2.summary().tables[0][1][1]))
for i in updated_equation.itertuples():
    print(' {:.3f} x ( {} ) '.format(float(str(i[2])),i[1]))
```

Logit: 0.95

Estimate [Logit] as L =

-0.014 x (Age)
-0.000 x (Income)
-0.373 x (Gender)
-0.456 x (VitD_levels)
-0.458 x (Doc_visits)
-0.190 x (vitD_supp)
+0.464 x (Initial_admin)
-0.402 x (Overweight)
-0.508 x (Arthritis)
-0.378 x (Anxiety)
-0.377 x (Allergic_rhinitis)
-0.374 x (Reflux_esophagitis)
-0.425 x (Asthma)
+0.404 x (Initial_days)
+0.183 x (Timely_admis)
-0.285 x (Timely_visits)
-0.500 x (Reliability)
-0.916 x (Options)
-0.317 x (Hrs_treat)
-0.349 x (Courteous)
-0.362 x (Active_listen)